

# 09. Processes

CAOS, MIPT 2019

# What is a process?

- An instance of the program in one of execution states
- Isolated virtual address space in UNIX

# Process attributes

- Memory
  - Registries, mmmaps, stacks
- Filesystem
  - Fds, pathes, umask
- Other
  - Envvars, limits, uids, ...

# Processes info

- `ps -A` - all process
- `top` – list all process by resources consumption
- `/proc` filesystem

# Tree of processes

- `Pid == 1` – `systemd` (init)
- Every process has a parent
- If parent dies before child, `systemd` becomes parent
- Parent needs to check if child dies

# Processes states

- Running
- Suspended – waiting for event, can be finished;
- sTopped – waiting for explicit wake up
- Zombie

# Schedulers

- FIFO
- Priority scheduling (man nice)
- Round-robin
- Multilevel FIFO
- [etc.](#)

# sched\_yield()

```
while (1) {  
    // do nothing - just waste CPU  
}
```

```
while (1) {  
    sched_yield(); // OK  
}
```



# Creating process

```
pid_t result = fork();
```

- Creates almost exact (except for ids, signals, timers, locks) copy of current process

- -1 – error
- 0 – parent for child process
- >0 – child for parent process

# Limits

- `/proc/sys/kernel/pid_max` – max number of simultaneously launched processes
- `/proc/sys/kernel/threads_max` – max number of simultaneously executing threads

# Fork bomb

```
void fork_bomb() {  
    pid_t p;  
    do {  
        p = fork();  
    } while (p != -1);  
    while (1) sched_yield();  
}
```

# Finishing process

- `_exit(int)` syscall
- `exit(int)` function (calls atexit handlers; flushes stdio; deletes tmpfiles)
- `return int` in main
- `kill -<signum> <pid>`

# Checking child

- `man wait / waitpid`
- `man wait3 / wait4`
- `WIFEXITED(wstatus)` – if was ended by `_exit`
- `WIFSIGNALED(wstatus)` – if was forced to die
- `WEXITSTATUS(wstatus)` – get exit code from 0 to 255
- `WTERMSIG(wstatus)` – get signal if was killed

```
int status;
waitpid(child, &status, 0);
if (WIFEXITED(status)) {
    printf("Exit code: %d",
WEXITSTATUS(status));
} else if (WIFSIGNALED(status)) {
    printf("Killed by %d signal",
WTERMSIG(status));
}
```

# Replace child's program

man 3 exec

- `l` – variadic arguments count; `v` – array of arguments
- `e` – pass envvars
- `p` – search programs in `$PATH`

```
int main() {  
    pid_t pid = fork();  
    if (pid == -1) { perror("fork"); exit(1);  
}  
  
    if (pid == 0) {  
        execlp("ls", "ls", "-l", NULL);  
        perror("exec"); exit(2);  
    } else {  
        waitpid(pid, NULL, 0);  
    }  
}
```



# Attributes of process that exec saves

- File descriptors
- Curdir
- Limits
- UID, GID